

Table of Contents

Introduction:.....	2
Abstract:	2
Objective:	3
Code & Output:	4
Importing necessary Libraries:	4
Dataset:.....	4
Dataset Preprocessing:	4
Outlier Analysis:	6
Time series:	7
Clustering:.....	9
Classification:.....	11
Conclusion:	12

Introduction:

In today's highly competitive business landscape, understanding and catering to customer needs has become a key priority for organizations. To achieve this, businesses need to delve deep into their customer data and extract meaningful insights. One powerful approach to accomplishing this is by applying clustering, time series analysis, and classification techniques to a dataset of customer invoices. These techniques enable businesses to identify customer segments, uncover trends in customer behavior, and predict customer churn. By harnessing the power of these analytical tools, businesses can enhance their marketing campaigns, optimize inventory management, and streamline staffing, leading to improved customer satisfaction and overall business performance.

Abstract:

This project explores the use of clustering, time series analysis, and classification on a dataset of customer invoices. These techniques can be used to group similar customers together, identify trends in customer behaviour, and classify customers based on their characteristics. This information can be used to improve marketing campaigns, prevent customer churn, and make better decisions about inventory, staffing, and marketing. This would allow you to target different marketing campaigns to different groups of customers. Time series analysis could be used to identify trends in customer behaviour. This would allow you to predict when customers are likely to buy and how much they are likely to spend. Classification could be used to predict whether a customer is likely to churn.

The following are benefits of using clustering, time series analysis, and classification on a dataset of customer invoices:

- Improved marketing campaigns: By understanding the different types of customers that you have, you can create marketing campaigns that are more likely to appeal to each group.
- Better inventory management: By understanding the demand for different products, you can ensure that you have the right amount of inventory on hand.
- More efficient staffing: By understanding the peak times of demand, you can ensure that you have the right number of staff on hand to meet customer needs.

Overall, the use of clustering, time series analysis, and classification on a dataset of customer invoices can provide a number of benefits for businesses. By understanding their customers better, businesses can create more effective

marketing campaigns, reduce customer churn, improve inventory management, and more.

Objective:

The primary objective of this project is to explore the potential of clustering, time series analysis, and classification in the context of customer invoice data. By analyzing this dataset, we aim to achieve the following goals:

- **Customer Segmentation:** Utilize clustering techniques to group customers with similar purchasing patterns, preferences, and characteristics. This segmentation will enable businesses to tailor their marketing efforts, providing personalized campaigns that resonate with each customer segment. By understanding the unique needs of different customer groups, businesses can maximize their marketing ROI and enhance customer engagement.
- **Trend Identification:** Apply time series analysis to uncover temporal patterns and trends in customer behavior. By identifying recurring patterns, seasonal fluctuations, and changes in purchasing habits, businesses can anticipate demand fluctuations, optimize inventory levels, and make informed decisions regarding production and supply chain management.
- **Enhanced Decision-Making:** By gaining a comprehensive understanding of customer segments, trends, and churn predictions, businesses can make data-driven decisions related to inventory management, staffing, and resource allocation. This will lead to optimized operational efficiency, cost savings, and improved customer satisfaction.

Code & Output:

Importing necessary Libraries:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from sklearn.impute import SimpleImputer
```

Dataset:

```
df = pd.read_excel("G:\DCS-SEM 4\Predictive analysis\Dataset\Superstore.xls")
df_cat = df['category']
df.head()
```

Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	Postal Code	Region	Product ID	Category	Sub-Category	Product Name	
0	1	CA-2016-152156	2016-11-08	2016-11-11	Second Class	CG-12520	Claire Gule	Consumer	United States	Henderson	42420	South	FUR-BO-10001798	Furniture	Bookcases	Bush Somerset Collection Bookcase
1	2	CA-2016-152156	2016-11-08	2016-11-11	Second Class	CG-12520	Claire Gule	Consumer	United States	Henderson	42420	South	FUR-CH-10000454	Furniture	Chairs	Hon Deluxe Fabric Upholstered Stacking Chairs,...
2	3	CA-2016-138688	2016-06-12	2016-06-16	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles	90036	West	OFF-LA-10000240	Office Supplies	Labels	Self-Adhesive Address Labels for Typewriters b...
3	4	US-2015-108966	2015-10-11	2015-10-18	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	33311	South	FUR-TA-10000577	Furniture	Tables	Bretford CR4500 Series Slim Rectangular Table
4	5	US-2015-108966	2015-10-11	2015-10-18	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	33311	South	OFF-ST-10000760	Office Supplies	Storage	Eldon Fold 'N Roll Cart System

Dataset Preprocessing:

```
In [3]: df.dtypes
```

```
Out[3]: Row ID          int64
Order ID         object
Order Date       datetime64[ns]
Ship Date        datetime64[ns]
Ship Mode        object
Customer ID       object
Customer Name     object
Segment          object
Country          object
City             object
State            object
Postal Code      int64
Region           object
Product ID       object
Category         object
Sub-Category     object
Product Name     object
Sales            float64
Quantity         int64
Discount         float64
Profit           float64
dtype: object
```

```
In [4]: df.shape
```

```
Out[4]: (9994, 21)
```

```
In [5]: obj = [x for x in df.columns if df[x].dtype == 'object']
```

```
In [6]: from sklearn import preprocessing
encoder = preprocessing.LabelEncoder()
for i in df.columns:
    if i in obj:
        df[i] = encoder.fit_transform(df[i])

df.head()
```

```
Out[6]:
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	Postal Code	Region	Product ID	Category	Sub-Category	Product Name	Sales	Quant
0	1	2500	2016-11-08	2016-11-11	2	143	166	0	0	194	42420	2	12	0	4	386	281.9600	
1	2	2500	2016-11-08	2016-11-11	2	143	166	0	0	194	42420	2	55	0	5	639	731.9400	
2	3	2296	2016-06-12	2016-06-16	2	237	201	1	0	266	90036	3	946	1	10	1433	14.6200	
3	4	4372	2015-10-11	2015-10-18	3	705	667	0	0	153	33311	2	318	0	16	366	957.5775	
4	5	4372	2015-10-11	2015-10-18	3	705	667	0	0	153	33311	2	1316	1	14	573	22.3880	

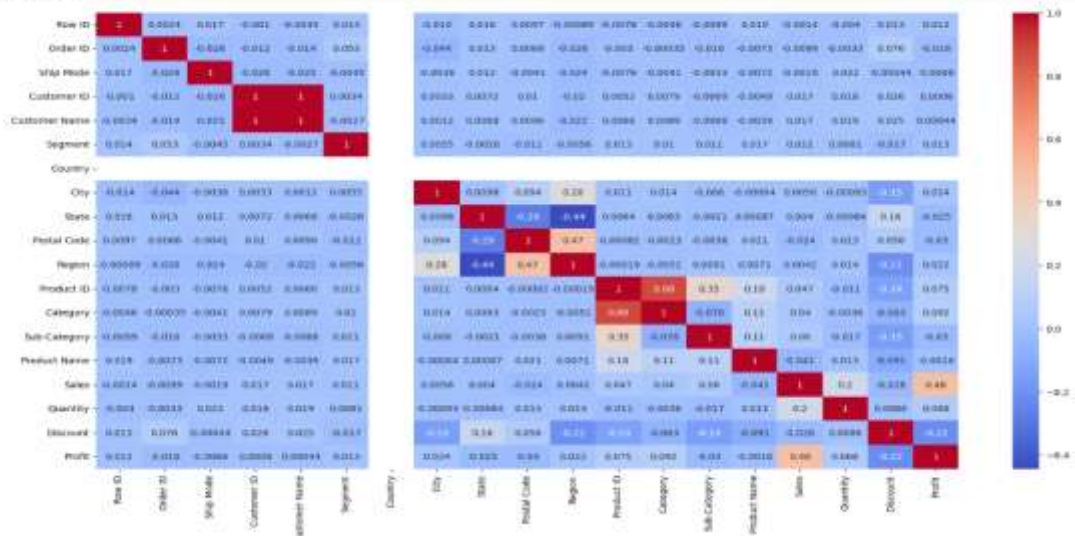
5 rows × 21 columns

```
In [7]: df.isnull().sum()
```

```
Out[7]: Row ID      0
Order ID      0
Order Date     0
Ship Date      0
Ship Mode      0
Customer ID     0
Customer Name   0
Segment        0
Country         0
City            0
State           0
Postal Code     0
Region          0
Product ID      0
Category        0
Sub-Category    0
Product Name    0
Sales           0
Quantity        0
Discount        0
Profit          0
dtype: int64
```

```
In [8]: corr = df.corr()
```

```
In [9]: plt.figure(figsize=(20,10))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.show()
```



Feature Selection

```
In [10]: drop = ['Row ID', 'Order ID', 'Ship Mode',
               'Customer Name', 'Segment', 'Country',
               'City', 'State', 'Postal Code', 'Region', 'Sub-Category']
df = df.drop(drop, axis = 1)
df.head()
```

```
Out[10]:
```

	Order Date	Ship Date	Customer ID	Product ID	Category	Product Name	Sales	Quantity	Discount	Profit
0	2016-11-08	2016-11-11	143	12	0	386	201.9600	2	0.00	41.9130
1	2016-11-08	2016-11-11	143	55	0	839	731.9400	3	0.00	219.5820
2	2016-06-12	2016-06-16	237	946	1	1433	14.6200	2	0.00	6.8714
3	2015-10-11	2015-10-18	705	319	0	365	957.5775	5	0.45	-383.0310
4	2015-10-11	2015-10-18	705	1316	1	573	22.3680	2	0.20	2.5164


```
In [11]: df.describe()
```

```
Out[11]:
```

	Customer ID	Product ID	Category	Product Name	Sales	Quantity	Discount	Profit
count	9994.000000	9994.000000	9994.000000	9994.000000	9994.000000	9994.000000	9994.000000	9994.000000
mean	400.480370	898.472383	0.972584	922.324795	229.858001	3.789674	0.156203	26.656896
std	228.585576	626.708445	0.629544	531.515675	623.245101	2.225110	0.206452	234.260188
min	0.000000	0.000000	0.000000	0.000000	0.444000	1.000000	0.000000	-6599.976000
25%	205.250000	453.000000	1.000000	474.250000	17.200000	2.000000	0.000000	1.726750
50%	405.500000	893.000000	1.000000	907.000000	54.490000	3.000000	0.200000	8.666500
75%	602.000000	1347.000000	1.000000	1390.000000	209.940000	5.000000	0.200000	29.364000
max	792.000000	1881.000000	2.000000	1849.000000	22638.480000	14.000000	0.800000	8399.976000

Outlier Analysis:


```
In [14]: sns.boxplot(df[["Sales", "Quantity", "Discount", "Profit"]])
plt.show()
```



```
In [15]: train = df[["Sales", "Quantity", "Discount", "Profit"]]
def impute_outliers(data, column, factor):
    q1 = data[column].quantile(0.25)
    q3 = data[column].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - factor * iqr
    upper_bound = q3 + factor * iqr

    data_copy = data.copy()
    data_copy[column] = np.where(data_copy[column] < lower_bound, np.nan, data_copy[column])
    data_copy[column] = np.where(data_copy[column] > upper_bound, np.nan, data_copy[column])
    imputer = SimpleImputer(strategy="mean")
    data_imputed = imputer.fit_transform(data_copy)

    return pd.DataFrame(data_imputed, columns=data.columns)
for column in ["Sales", "Quantity", "Discount", "Profit"]:
    train = impute_outliers(train, column, 1.5)
df[["Sales", "Quantity", "Discount", "Profit"]] = train
sns.boxplot(train)
plt.show()
```



Time series:

ARIMA Forecasting

```
In [17]: from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error, mean_squared_error
import warnings
warnings.filterwarnings("ignore")
```

Model Fitting and evaluation

```
In [18]: arima_model = ARIMA(sales_by_year["Sales"], order=(1, 1, 1))
arima_model_fit = arima_model.fit()
```

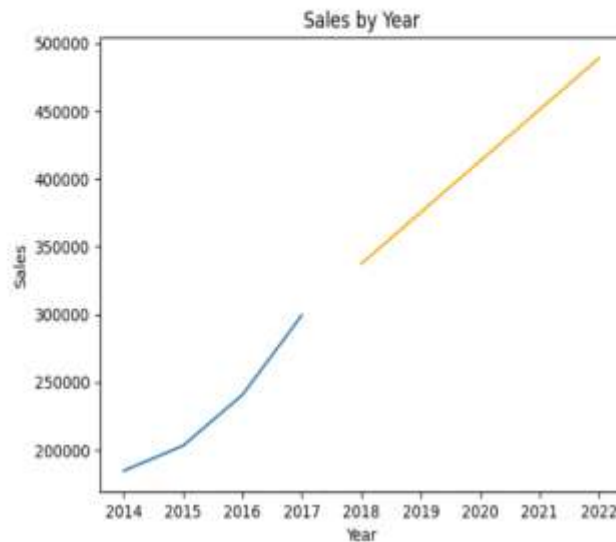
Forecasting

```
In [19]: forecast = arima_model_fit.forecast(steps = 5).round(2)
print(forecast)

4    337297.66
5    375287.18
6    413114.25
7    451019.13
8    488921.73
Name: predicted_mean, dtype: float64
```

Plotting of Actual data with Forecasted data

```
In [20]: plt.plot(sales_by_year["Order Year"], sales_by_year["Sales"])
plt.plot([2018, 2019, 2020, 2021, 2022], forecast, color = "orange")
plt.xlabel("Year")
plt.ylabel("Sales")
plt.title("Sales by Year")
plt.show()
```



Forecasting for each Category

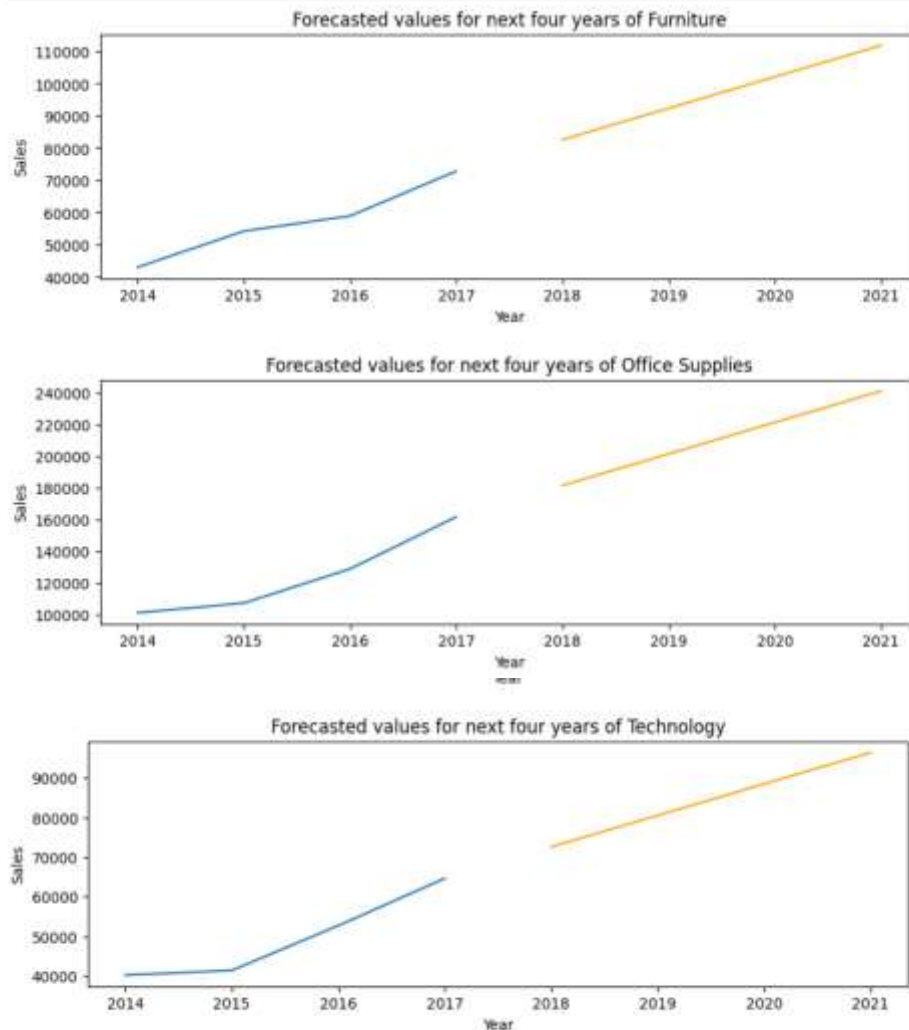
```
In [21]: x = df.groupby(['Category', 'Order Year'])['Sales'].sum().reset_index().pivot(index = 'Order Year',
                                                                                      columns = 'Category', values = "Sales")
x.head()
```

```
Out[21]:
```

	Category	Furniture	Office Supplies	Technology
Order Year				
2014	43028.096412	101308.051450	40193.214155	
2015	54214.385590	107482.871647	41380.777012	
2016	58929.683552	129057.724561	52706.806171	
2017	72848.195671	161924.363365	64815.391020	

```
In [22]: def forecasting(x):
    arima_model = ARIMA(x, order = (1,1,1))
    arima_model_fit = arima_model.fit()
    forecast_values = arima_model_fit.forecast(steps = 4)
    return forecast_values.round(2)

for i in x.columns:
    values = forecasting(x[i])
    plt.figure(figsize = (10, 3))
    plt.plot(x.index, x[i])
    plt.plot([2018, 2019, 2020, 2021], values, color = 'orange')
    plt.xlabel('Year')
    plt.ylabel('Sales')
    plt.title(f'Forecasted values for next four years of {i}')
```

Clustering:

Clustering

```
In [23]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler
```

```
In [24]: df_cluster = df.groupby("Customer ID").sum()
train = df_cluster[["Sales", "Quantity", "Discount", "Profit"]]
```

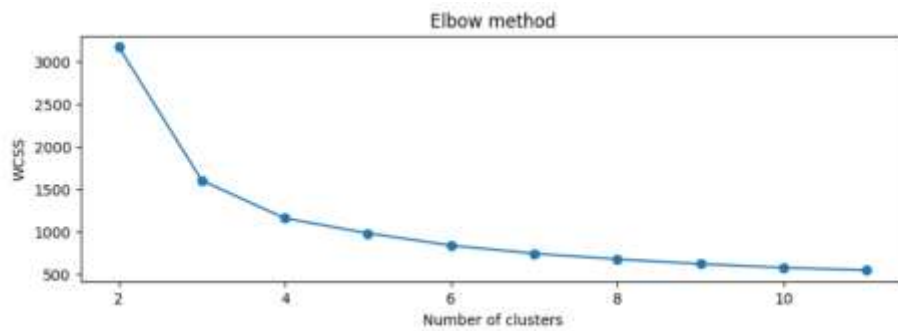
Optimal K size

Elbow Method

```
In [25]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(train)

#Elbow Method
wcss = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init = 10)
    kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_)

plt.figure(figsize = (10, 3))
plt.plot(range(2, 12), wcss, marker="o")
plt.xlabel("Number of clusters")
plt.ylabel("wcss")
plt.title("Elbow method")
plt.show()
```



Inference:

We can see that from 2 - 3 the value of WCSS decreasing fast and in a straight but after 2 there is a slow drop and decreasing exponentially so $K = 3$

K - Means Algorithm

```
In [26]: k_opt = 3
kmeans_opt = KMeans(n_clusters=k_opt, random_state=42, n_init = 10)
kmeans_opt.fit(train)

# Assign the cluster labels to the data set
df_cluster["cluster"] = kmeans_opt.labels_
df_cluster.head()
```

```
Out[26]:
```


	Product ID	Product Name	Sales	Quantity	Discount	Profit	Order Month	Order Year	Cluster
Customer ID									
0	10075	8722	1145.625495	35.000000	0.903468	121.282405	58	22199	0
1	19095	13133	914.829495	45.000000	1.000000	217.682303	128	30237	1
2	12695	14127	1790.632000	36.000000	0.200000	195.867205	84	24191	0
3	13339	17334	1488.231473	70.654077	1.600000	240.832309	162	36274	0
4	5041	5259	831.254000	21.000000	0.103468	82.141903	33	12088	1

```
In [27]: df_cluster[["Sales", "Quantity", "Discount", "Profit", "cluster"]].groupby("cluster").mean(numeric_only = True).round(2)
```

```
Out[27]:
```

	Sales	Quantity	Discount	Profit
cluster				
0	1367.79	81.86	1.50	164.61
1	594.46	29.44	0.83	87.28
2	2399.70	81.18	2.20	277.66

```
In [28]: import matplotlib.pyplot as plt
plt.scatter(df_cluster['Sales'], df_cluster['Profit'], c=df_cluster['cluster'])
plt.xlabel('Sales')
plt.ylabel('Profit')
plt.show()
```



Inference:

- Cluster 0 has an average sales value of 1367.79 , an average quantity of 51.86, an average discount of 1.50 and an average profit of 164.51.
- Cluster 1 has an average sales value of 594.49, an average quantity of 29.44 , an average discount of 29.44 and an average profit of 87.26.
- Cluster 2 has an average sales value of 2399.70 , an average quantity of 81.18, an average discount of 2.20 and an average profit of 277.66.

Conclusion:

Based on these values, Cluster 2 has the highest mean values for sales, quantity, discount and profit. This suggests that customers in Cluster 2 are generating more sales and profit for the supermarket compared to customers in Clusters 0 and 1.

Overall, this analysis suggests that customers in Cluster 2 may be more valuable to the supermarket in terms of generating sales and profit. It may be worth exploring further to understand what differentiates customers in Cluster 0 from those in Clusters 0 and 1 and how the supermarket can attract more customers like those in Cluster 2.

Classification:

Classification

```
In [29]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
```

Fitting and Evaluation

```
In [30]: X = df_cluster[["Sales", "Quantity", "Discount", "Profit"]]
y = df_cluster["cluster"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
log_reg = LogisticRegression(max_iter = 150)
log_reg.fit(X_train, y_train)

#Predictions
y_pred = log_reg.predict(X_test)

#Evaluation
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[57  2  3]
 [ 0 73  0]
 [ 4  0 20]]
```

	precision	recall	f1-score	support
0	0.93	0.92	0.93	62
1	0.97	1.00	0.99	73
2	0.07	0.83	0.85	24
accuracy			0.94	159
macro avg	0.93	0.92	0.92	159
weighted avg	0.94	0.94	0.94	159

Classifying new rows

```
In [31]: def predict(test):
        test = scaler.transform(test)
        return log_reg.predict(test)

name = input('Name:')
sales = int(input("Sales:"))
quantity = int(input("Quantity:"))
discount_rate = float(input("Discount:"))
profit = int(input("Profit:"))

cluster = predict(pd.DataFrame({'Sales': [sales], 'Quantity': [quantity], 'Discount': [discount_rate], 'Profit': [profit]}))
print(f'Customer {name} belongs to cluster :{cluster[0]}')

Name:Harish
Sales:180000000000
Quantity:2123
Discount:0.1
Profit:100020
Customer Harish belongs to cluster :2
```

Conclusion:

The value of clustering, time series analysis, and classification techniques in extracting actionable insights from customer invoice data. The project's findings provide businesses with the knowledge and tools to improve marketing effectiveness, optimize inventory management, reduce customer churn, and make informed decisions for sustainable growth and success in the market.